CSE 415 Autumn 2019   Assignment 4

Last name: _Tianyuan_   First name: _Fu_

Due Thursday night October 24 via Gradescope at 11:59 PM. A maximum of one late day may be used on this assignment, with the usual 10 percent reduction in score.

Do the following three exercises. These are intended to take 20-30 minutes each if you know how to do them. Each is worth 30 or 35 points.

---

1   Efficiency of Search Algorithms

Suppose a binary tree-searching algorithm needs to search to a maximum depth of 4. However, it could find a goal node and stop there, at any level d of the tree, $0 \le d \le 4$.

(a) (3 points) How many node visitations in the tree above must be done by each of the following algorithms, to find the goal node marked G? (Assume children are handled in left-to-right order.)

DFS: 22   BFS: 13   IDDFS: 23

(b) (9 points) Now for a complete binary tree having N vertices, give Big-Theta characterizations of the worst-case running times:

DFS: $\Theta(N)$ BFS: $\Theta(N)$ IDDFS: $\Theta(N)$

1

(c) (6 points) Now for a general tree having N vertices, and any branching factor, including 1, or even irregular (varying numbers of children, but obviously always non-negative numbers), give Big-Theta characterizations of the worst-case running times:

DFS: $\Theta(N)$ BFS: $\Theta(N)$ IDDFS: $\Theta(N^2)$

(d) (6 points) Using Big-Theta notation again, describe the amount of memory required by each algorithm for the general tree of N vertices.

DFS: $\Theta(\log N)$ BFS: $\Theta(N)$ IDDFS: $\Theta(\log N)$

(e) (4 points) Suppose we now want to use graph-search versions of these algorithms. What is an important benefit of BFS and IDDFS that DFS does not have?

The result of BFS and IDDFS are assured to be the optimized solution, while DFS may reach the goal state with a non-optimized route.

(f) (3 points) What impact does a heuristic function have on the run time of the A* algorithm? Assume that the heuristic is admissible and that we don't consider the cost of computing the heuristic function itself.

The usage of heuristic functions can reasonably reduce the run time of the A* algorithm. Heuristic function will give an estimation of cost from the current state to the goal state, such that the algorithm can select the next state with lowest cost (highest priority), rather than blindly choose from left to right or right to left.

(g) (4 points) What does an admissible heuristic do for us in the best and worst cases? (Consider the cost of computing the heuristic function as well as other issues.)

Best case: heuristic function is easy to compute and the value of it can always represent the real cost from the current state to the goal state. The algorithm using this function would have low overhead and will reach the goal state with the lowest cost route.

Worst case: heuristic function requires huge computation power and results have a large deviation from the actual cost. It may lead the algorithm to choose a non-optimized route or unable to reach the goal state.

2 Heuristic Search

Imagine that you are part of a game-design team creating a game that takes place in a hexagonal grid world. Your job is to evaluate some heuristics proposed to help the monsters of the game do path planning.

The figure below shows a general hexagonal grid in which each cell is identifiable by a triple of coordinates: x (which tells how far it is from the origin on an axis that makes an angle of $\pi/6$ with the horizontal axis), y (which tells how far from the origin it is on an axis that makes an angle of $5\pi/6$ with the horizontal axis), and z (which tells how far down from the origin it is on the vertical axis). Note that there is (intentionally) some redundancy here. In particular, $x + y + z = 0$.

We define the distance between two adjacent cells as follows:

$$d((x_1, y_1, z_1), (x_2, y_2, z_2)) = |x_1 - x_2| + 2 \cdot |y_1 - y_2| + 4 \cdot |z_1 - z_2|$$

The monsters in the game need to perform path planning for two reasons. One is to escape from fighting with the protagonist. The other is to move towards the protagonist to make an attack. (Although the monsters should not be too intelligent, we will assume that they do need to know what the lowest-cost path is between their starting states and destination states.)

Suppose the operators are the following, in the sequence given. Note that only the state-transformation portion of each operator is shown here, but they all have a precondition that monsters are not allowed to go out of the grid or walk through walls. Also note that the order in which the operators are listed has significance during the various searches.

$$\phi_0: \; x \leftarrow x + 1 \text{ and } y \leftarrow y - 1. \quad \overset{3}{\underset{}{\longrightarrow}}$$

$$\phi_1: \; z \leftarrow z + 1 \text{ and } y \leftarrow y - 1. \quad \overset{6}{\searrow}$$

$$\phi_2: \; z \leftarrow z + 1 \text{ and } x \leftarrow x - 1. \quad \overset{5}{\swarrow}$$

$$\phi_3: \; y \leftarrow y + 1 \text{ and } x \leftarrow x - 1. \quad \overset{3}{\longleftarrow}$$

$$\phi_4: \; y \leftarrow y + 1 \text{ and } z \leftarrow z - 1. \quad \overset{6}{\nwarrow}$$

$$\phi_5: \; x \leftarrow x + 1 \text{ and } z \leftarrow z - 1. \quad \overset{5}{\nearrow}$$

The cost of going from one cell to a neighbor depend on the operator. For example, applying ϕ_3 to state M leads to state $(-1, +1, 0)$, so that the distance value, by the formula above, is 3. The cost of a path is, as usual, the sum of its edge costs.

(a) (6 points) Consider the problem of finding a lowest-cost path from M to P in the example game situation below. The thick black walls in the diagram represent barriers through which the monsters must not pass. (Therefore, the precondition for operator ϕ_0 is not satisfied at state $(+1, -1, 0)$, and neither is the precondition for operator ϕ_1, so that this state has only 4 successors.)

What is the cost of the shortest path? (i.e., total distance for all its edges)

$6 + 3 + 3 + 5 = 17 \qquad \phi_1 \Rightarrow \phi_0 \Rightarrow \phi_0 \Rightarrow \phi_5.$

(b) (6 points) What is the number of states that would be expanded by UCS to find a lowest-cost path from M to P?

19.

(c) (4 points) Consider the heuristic $h_x((x, y, z)) = \min(|x - x_g|, |y - y_g|, |z - z_g|)$, where (x_g, y_g, z_g) refer to the coordinates of the goal cell.

Determine whether or not h_x is admissible and explain why it is or why it is not.

hx is admissible since:
The route from the start state to the goal state in the diagram must be longer or equal to the projected distance to any of the three dimensions.
All of the three elements in the min function are bigger or equal to 0
Without any barrier, any of the cell in the figure can have a path to the final cell.

(d) (6 points) Determine how many states would be expanded by A* using h_x to find a shortest path from M to P.

15

(e) (5 points) Explain how the use of the heuristic is affecting the search (compared with UCS).

The usage of the heuristic function can reduce the number of expended states, since it will guide the algorithm to choose the cells that moves towards the final state, rather than in the opposite direction.

(f) (5 points) Propose another heuristic that you believe will outperform h_x. (Do not propose the exact distance defined above or function that includes it.) Give a formula to define your function. Also give an argument for or against its admissibility. Why do you believe it will outperform h_x?

$$h_x = \frac{1}{2}\left|\,(\,|x_g| + |y_g| + |z_g|\,) - (\,|x| + |y| + |z|\,)\,\right|.$$

This function must be admissible, for it represents the difference of distances to the enter of the fissure. It must be shorter than the actual path.

(g) (3 points) How many states will be expanded by A* using your heuristic to find a shortest path from M to P?

16

3 Markov Decision Processes

Consider a situation in which a robot monkey is in a room in which some bananas are hanging from the ceiling. The monkey can move its left arm to any one of 3 different places, and it might be holding or not holding bananas in either of two states. The monkey starts out not holding bananas with left hand low (state F). It can try to raise its hand, and so it can transition to state A. However, this monkey has a somewhat unreliable body, and most actions are only effective with probability 0.7. (Then, with probability 0.3, the arm goes a different way. But there is 0 probability of staying in the same state.) The action of raising the hand from state F, however, is the one completely reliable and deterministic action, so the next state after that action is guaranteed to be state A. If the monkey again tries to raise his hand from state A, there is a 0.7 probability of arriving in state B, which is roughly where the bananas are. There is a probability of 0.3 that the action will cause the monkeys hand to move back down to state F. From state B, the monkey has a choice of action: (G) grab the bananas and go down to state C, or (D) leave the bananas and go down to state A. (These again are nondeterministic, so action G in state B goes to C with probability 0.7 and to state A with probability 0.3; action D in state B goes to A with probability 0.7 and to C with probability 0.3.) From state C, the monkey can go to state E or back up to state B. From state E, the monkey can finally get the bananas all the way down and consume them, for a reward of +10. Note that the other rewards are either effort-consuming (raising the hand or raising it and the bananas for rewards of -0.2 or -0.5) or somewhat easier, lowering the hand (for rewards 0f 0.1).

(a) (10 points) Create a table representing the function $T(s, a, s')$ for this MDP. Assume that there are three possible actions: U, D and G. U means try to go up, and D means try to go down. G means try to grab the bananas (and go down from B to C). Action U is applicable only at states A, C, E, and F. Action D is applicable only at states A, B, C, and E. Action G is only applicable at state B. In this table you do not have to provide an explicit entry for any $T(s, a, s')$ whose value is 0. Hint: there should be 17 entries in your table.

$s \in S$	$a \in A$	$s' \in S'$	$T(s,a,s')$
A	u	B	0.7
A	u	F	0.3
A	D	F	0.7
A	D	B	0.3
B	G	C	0.7
B	G	A	0.3
B	D	A	0.7
B	D	C	0.3
C	u	B	0.7
C	u	E	0.3
C	D	E	0.7
C	D	B	0.3
E	u	C	0.7
E	u	F	0.3
E	D	F	0.7
E	D	C	0.3
F	u	A	1

(b) (5 points) Create a table representing the function $R(s, a, s')$. There should be one explicit entry here for each explicit entry you gave in part (a).

$s \in S$	$a \in A$	$s' \in S'$	$T(s,a,s')$	$R(s,a,s')$
A	u	B	0.7	-0.2
A	u	F	0.3	+0.1
A	D	F	0.7	+0.1
A	D	B	0.3	-0.2
B	G	C	0.7	+0.1
B	G	A	0.3	+0.1
B	D	A	0.7	+0.1
B	D	C	0.3	+0.1
C	u	B	0.7	-0.5
C	u	E	0.3	+0.1
C	D	E	0.7	+0.1
C	D	B	0.3	-0.5
E	u	C	0.7	-0.5
E	u	F	0.3	+10
E	D	F	0.7	+10
E	D	C	0.3	-0.5
F	u	A	1	-0.2

(c) (10 points) Create a table representing the expected reward for taking action a in state s. Call this function $Q_1(s, a)$. Note: if action a is not allowed in state s then $S_1(s, a) = 0$. The expected reward is the *expectation* of the reward from state s given that the action

is s. Because of the nondeterminism in the actions, this is normally a weighted sum of the possible rewards out of s, and the weight corresponding to action a is usually the largest of the weights. Hint: there should be 9 entries in your table.

$s \in S$	$a \in A$	$Q_1(s, a)$
A	U	-0.11
A	D	$+0.01$
B	G	$+0.1$
B	D	$+0.1$
C	U	-0.32
C	D	-0.08
E	U	$+2.65$
E	D	$+6.85.$
F	U	-0.2

(d) (5 points) What is the total reward corresponding to the following episode, where the monkey starts in state A and makes the transitions below?

A-D-F-U-A-U-B-D-A-U-B-G-C-D-B-G-C-D-E-D-C-D-E-D-F-U-A

(Add up all the actual rewards here; do not use expected rewards.)

8.6